THCHS-30 : A Free Chinese Speech Corpus

Article · December 2015

Source: arXiv

2 authors, including:

Dong Wang
The University of Edinburgh
139 PUBLICATIONS   **823** CITATIONS

THCHS-30 : A Free Chinese Speech Corpus

Dong Wang[*] and Xuewei Zhang

Abstract

Speech data is crucially important for speech recognition research. There are quite some speech databases that can be purchased at prices that are reasonable for most research institutes. However, for young people who just start research activities or those who just gain initial interest in this direction, the cost for data is still an annoying barrier. We support the 'free data' movement in speech recognition: research institutes (particularly supported by public funds) publish their data freely so that new researchers can obtain sufficient data to kick off their career. In this paper, we follow this trend and release a free Chinese speech database THCHS-30 that can be used to build a full-fledged Chinese speech recognition system. We report the baseline system established with this database, including the performance under highly noisy conditions.

Keywords: TCMSD database; THCHS-30; speech recognition; deep neural network

1 Introduction

1.1 Speech recognition and data

Automatic speech recognition (ASR) has been an active research area for several decades. Traditional ASR systems are based on the hidden Markov model (HMM) to represent temporal dynamics of speech signals and the Gaussian mixture model (GMM) to represent distributions of the signals within a stationary short period of time that generally corresponds to a pronunciation unit, e.g., a phone. This HMM-GMM paradigm dominates ASR research for nearly three decades, until a few years ago when deep learning stepped in and attained revolutionary success. Although several variants have been proposed, the most popular deep learning approach for ASR is the so called HMM-DNN hybrid architecture, where the HMM framework is retained while the GMM is substituted for a deep neural network (DNN) to model static (stationary) distributions of speech signals. Readers can refer to [1, 2, 3, 4, 5, 6] for details about DNN-based ASR methods.

Due to the great improvement caused by DNN and other factors (e.g., mobile network and big data), speech recognition has gained much attention in the research community and multiple industrial sectors. On one hand, the greatly improved accuracy (in terms of word error rate, WER) has enabled broad practical applications, e.g., voice search on mobile devices; on the other hand, the technique is still fragile and it is not easy to recognize speech with strong accents and in complex noise conditions. It is also not simple to migrate a system that works well in one domain to another if the difference is significant. This subtle situation attracts much interest on speech technologies among researchers, both young graduates and senior researchers who previously work on different areas. This is certainly good news since more researchers mean more rapid advance in this field, and more close we are from the ultimate goal of understanding human speech.

Nevertheless, there is a big barrier that is hindering researchers joining the ASR community: data. Everybody knows that ASR research heavily relies on data, even more heavily than other areas that look similar, e.g., face recognition. This is because human speech is so complex that the essential voice patterns can be learned only with a large amount of data.

1.2 Data is blocking research

There are indeed quite some speech databases available for ASR research, e.g., the TIMIT database [7] that was designed for small-scale phone recognition, the WSJ database [8] that was designed for large-scale continuous speech recognition on the domain of broadcast news, the Switchboard database [9] that was designed for large-scale recognition for telephone conversions, and the AMI database [10] that consists of meeting recordings. These databases can be purchased from the institutes that created the data, but more generally from commercial organizations including LDC[1] and ELRA[2]. For Chinese, the most popular database is the RAS 863 corpus [11], which involves continuous reading speech of more than 80 speakers, resulting in

[*]Correspondence: wangdong99@mails.tsinghua.edu.cn
Center for Speech and Language Technology, Research Institute of Information Technology, Tsinghua University, ROOM 1-303, BLDG FIT, 100084 Beijing, China
Full list of author information is available at the end of the article

[1]https://www.ldc.upenn.edu/
[2]http://catalog.elra.info/

nearly 100 hours of speech signals. Other commercial speech databases can be purchased from Datatang[3] and Speech Oceanf[4]. The advantage of using these databases is that they are standard and publicly available, so that results from different authors become comparable. This is very important for researchers to identify true promising algorithms and is actually one of the most important factors that contribute to the recent success of the ASR research.

However, almost all these databases are not free; in fact, most of them are quite expensive. For example, the middle-sized WSJ corpus that was recorded 20 years ago asks for $2,500$ US dollars[5]. This licence fee is affordable for most research institutes, however for individual researchers who just start their research activities or those who foster initial interests on ASR, $2,500\$$ is rather expensive, particularly for researchers in developing countries. We can not complain the high licence fee because organizing the recordings and annotation is very costly, and so it is not feasible to expect commercial institutes to provide free data. Nevertheless, the high licence fee does form a barrier that blocks the initial interest to ASR of beginners. This is an intolerable situation in the era where research activities have become more and more open, collaborative, and sharing.

1.3 Free data movement
We advocate a 'free data' movement in the speech community. Data should not be private properties if they are collected by public institutes such as universities and public-funded research groups. Since the collection and annotation is supported by public funds, the data should be free for people who pay the fund, or even for more broad people if the freedom will result in more benefit for the people who originally pay for the research. It is a shame that research projects supported by public funds use money from taxpayers to build up private resources. Additionally, even commercial companies should consider release free data, because the freedom will ultimately benefit the industrial sector: free data will invoke more research in this area and so more opportunities to enjoy novel technic development and professional employees. Finally, the development of internet accumulates large amount of speech data and the development of ASR technics allows cheaper annotation for these data. This makes the cost of data collection is not as expensive as before, and thus free data is more feasible.

There have been some initial attempts towards this direction. For example, the EU-supported AMI/AMID

project released all the data (both speech and video recordings on meetings), the VoxForge project[6] provides a platform to publish GPL-based annotated speech audio. OpenSLR is another platform to publish open resources for speech and language research, including the LibriSpeech[7]. The Sprakbanken database[8] is another free database in Swedish, Norwegian, Danish. For Chinese, we have never seen a free speech database that is sufficient enough to build a full-fledged large-vocabulary Chinese ASR system.

1.4 Our release
To respect the creed that 'public research should publish data' and break the data monopoly, to protect the initial interest of individual researchers towards this research direction, the center for speech and language technologies (CSLT) at Tsinghua University recently started publishing free databases. The first database we published is THUYG-20, which is used for training Uyghur speech recognition systems [12]. This publication is collaborated with the Xinjiang university and the Tsinghua AI Cloud Research Center (AICRC)[9].

In this paper, we continue the free speech data process. We publish a Chinese database that was recorded by the first author 15 years ago [13]. This database consists of 35 hours of speech signals recorded from 50 participants. Associated with the speech signals, we provide the full set of resources including lexicon, LM and the training recipe, so that new practitioners can use these resources to establish a basic large vocabulary continuous Chinese speech recognition system. Interested readers can download the data freely from http://data.cslt.org/thchs30/README.html. To our best knowledge, this is the first release that can be used to build a practical Chinese speech recognition system. We hope this release can provide a standard reference for Chinese ASR researchers and invoke much enthusiasm of young people.

The rest of the paper is organized as follows: Section 2 overviews the features of the released database, which we refer to as THCHS-30, and Section 3 describes the baseline system we built based on the release data and resource. The paper is concluded in Section 4.

2 Features of THCHS-30
We describe the THCHS-30 speech database in this section. This database was recorded in 2000 - 2001 by

[3]http://www.datatang.com/

[4]http://www.speechocean.com/

[5]https://catalog.ldc.upenn.edu/LDC94S13A

[6]http://www.voxforge.org/

[7]http://www.openslr.org/12/

[8]http://www.nb.no/sbfil/talegjenkjenning/

[9]http://data.cslt.org/thuyg20/README.html

the first author when he was a master student, supervised by Prof. Xiaoyan Zhu [13]. The design goal was a supplemental data resource for the 863 database, by which the phone coverage can be maximally improved. The new database was named as TCMSD (Tsinghua Continuous Mandarin Speech Database), and has remains private usage since its creation. We publish the data 15 years later, with the permission from the licence holder, Prof. Zhu. The database was renamed as THCHS-30, standing for 'Tsinghua Chinese 30 hour database'. The same nomenclature has been used for the THUYG-20 database, and will be followed by a number of free databases that we will publish soon.

2.1 Speech signals

THCHS-30 involves more than 30 hours of speech signals recorded by a single carbon microphone at the condition of silent office. Most of the participants are young colleague students, and all are fluent in standard Mandarin. The sampling rate of the recording is 16,000 Hz, and the sample size is 16 bits.

The sentences (text prompts in recording) of THCHS-30 were selected from a large volume of news wise, with the goal of augmenting the 863 database with more phone coverage. We selected 1000 sentences for recording. Table 1 reports the bi-phone and tri-phone coverage with/without THCHS-30 (reproduced from [13]). It can be seen that THCHS-30 indeed improves phone coverage of the 863 database.

Table 1: Phone coverage with THCHS-30

Database	863	THCHS-30	863 + THCHS-30
No. Sentence	1500	1000	2500
Bi-phone Coverage	58.4%	71.5%	73.4%
Tri-phone coverage	7.1%	14.3%	16.8%

The recordings are split into four groups according to the recording text: A (sentence ID from 1 to 250), B (sentence ID from 251 to 500), C (sentence ID form 501 to 750), D (sentence ID from 751 to 1000). The utterances of group A, B and C are combined as the training set, which involves 30 speakers and 10893 utterances. All the utterances in group D are used as the test set, which involves 10 speakers and 2496 utterances. The statistics of THCHS-30 are represented in Table 2.

Table 2: Statistics of THCHS-30 database

Data Set	Speaker	Male	Female	Age	Utterance	Duration (hour)
Training	30	8	22	20-55	10893	27.23h
Test	10	1	9	19-50	2496	6.24h

2.2 Additional resources

To assist constructing a practical Chinese ASR system, some additional resources are published in the THCHS-30 release. These resources include lexica, language models, training recipes and some useful tools. Additional data in noise conditions are also provided.

2.2.1 Lexicon and LM

We release two language models and the associated lexica. The word-based LM involves 48k words and is based on word 3-grams, and the phone-based LM involves 218 Chinese tonal finial-initials and is based on phone 3-grams. The word LM was trained using a text collection that was randomly selected from the Chinese Gigaword corpus[10]. The training text involves 772,000 sentences, amounting to 18 million words and 115 million Chinese characters. The phone LM was trained using a much smaller text collection involving 2 million characters. The reason to use a smaller text data for phone LM training is that we want to retain linguistic information as little as possible in the model so that the resultant performance reflects more directly the quality of acoustic models. The two LMs were trained with the SRILM tool [14].

2.2.2 Scripts and recipe

The recipe and some utility scripts are also published for training a complete Chinese ASR system with THCHS-30. The scripts are based on the framework of the Kaldi toolkit [15]. The training process is basically similar to the wsj s5 GPU recipe provided by Kaldi, although some modifications have been made to make it suitable for Chinese ASR.

2.2.3 Noise data

We are also interested in ASR tasks in noisy conditions, therefore provide a noisy version of THCHS-30: all the training and test data were corrupted by three types of noise (separately): white noise, car noise and cafeteria noise. We focus on the 0 db condition, which means that energy of noise and speech signals are the same (thus very noisy). The noise embedding was performed by simple wave mixture, where the noise signals are obtained from the public DEMAND noise repository[11].

2.3 Call for challenge

Since THCHS-30 is public and free, everyone can download it, build their systems and compare their results with others. To encourage the research, we call for challenge based on the provided resources, including two tasks: large vocabulary recognition and phone recognition. Although the former is more close to practical usage, the latter can test acoustic modeling approaches in a more focused way. For each task, we also compete for performance on the 0 db noise condition with the three types of noise: white, car and cafeteria.

[10]https://catalog.ldc.upenn.edu/LDC2003T09
[11]http://parole.loria.fr/DEMAND/

3 Baseline system

We describe our baseline system built with THCHS-30, and report the performance in various conditions that we have called for challenge. We treat these results as the baseline for the competition; any improvement by any authors will be published on the competition web page[12].
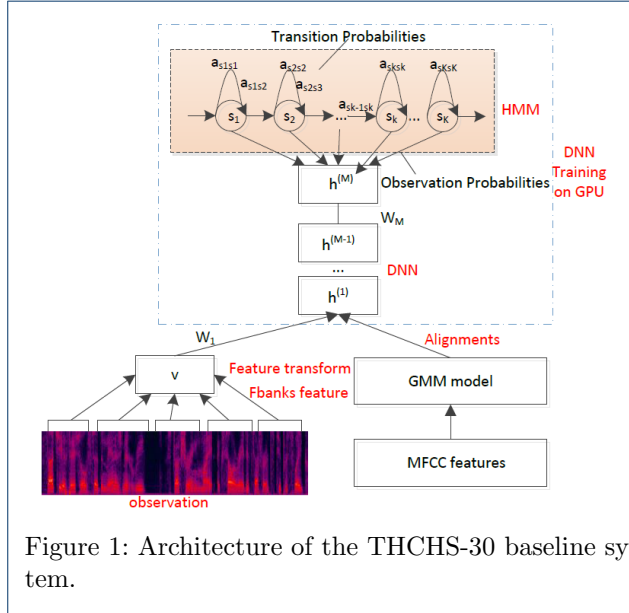
3.1 Framework and setting

Figure 1: Architecture of the THCHS-30 baseline system.

We use the Kaldi toolkit to train the HMM-DNN hybrid acoustic models. The architecture of the system is shown in Fig 1. Following this architecture, a monophone GMM system was firstly trained with the standard 13 dimensional MFCCs plus the first and second order derivatives. The cepstral mean normalization (CMN) was employed to reduce the channel effect. A triphone GMM system was then constructed based on the monophone system with features transformed by LDA and MLLT. The final GMM system was used to generate state alignment for the subsequent DNN training [16].

The DNN system was trained based on the alignments provided by the GMM system. The features were 40-dimensional FBanks, and neighbouring frames were concatenated by a 11-frame window (5 windows on each side). The concatenated features were transformed by LDA, by which the dimensionality was reduced to 200. A global mean and variance normalization was then applied to obtain the DNN input. The DNN architecture consists of 4 hidden layers and each layer consists of 1200 units. The output layer

[12]http://data.cslt.org/thchs30/challenges/asr.html

is composed of 3386 units. The baseline DNN model was trained with the criterion of cross entropy. The stochastic gradient descendent (SGD) algorithm was used to performance the optimization. The mini batch size was set to 256 frames, and the initial learning rate was set to be 0.008.

3.2 Preliminary results

The performance of the DNN system based on clean training data is presented in Table 3, where the word ASR task is evaluated in terms of character error rate (CER) and the phone ASR task is evaluated in terms of phone error rate (PER). Table 3 reports the test results on both the clean speech and the 0 db noisy speech. It can be seen that the performance on noisy speech is much lower than that on clean speech, particularly when the corruption is white noise. Note that the CER and PER are rather low compared to those obtained with many standard databases (e.g., 863 database), even on clean speech. This is mostly attributed to the 'abnormality' of the test data: to seek for maximal phone coverage, the sentences selected for THCHS-30 tend to be rather peculiar in both pronunciation and spelling. This makes the recognition task highly challenging, and solving the peculiarity perhaps needs some specific techniques for example more aggressive phone sharing.

Table 3: Performance of THCHS-30 baseline system trained with clean speech.

Condition	Word Task (CER%)	Phone Task (PER%)
GMM Clean	50.71	30.53
DNN Clean	43.19	25.16
0db-White	96.89	71.79
0db-Car	46.76	27.88
0db-Cafeteria	82.57	56.71

3.3 Noise cancellation with DAE

The significant performance degradation on noisy data can be addressed in multiple ways, for example, multiple condition training [17] and noisy training [18, 19]. Although promising, these approaches require re-training the DNN model and so is not convenient for dealing with new noise types. This paper applies a noise cancellation approach based on deep autoencoder (DAE). DAE is a special implementation of autoencoder (AE), by introducing random corruptions to the input features in model training. It has been shown that this model is very powerful in learning low-dimensional representations and can be used to recover noise-corrupted input [20]. In [21], DAE is extended to a deep recurrent structure and has been employed to recover clean speech in noisy conditions for ASR. A recent study employs DAE in de-reverberation [22], and

our work also showed that it can be used for removing music corruption [23].

We employ a simple DAE for noise cancellation in this study. Firstly noisy data are generated by corrupting clean speech with noise signals, and then a non-linear transform function can be learned with these data. Since the noisy data are generated artificially, the only cost is to find a piece of noise signal and use it to corrupt clean speech. The noisy data generation process follows a scheme similar to the one in [19], where a particular SNR is sampled for each clean utterance, following a Gaussian distribution whose mean is 0 db. The sampled SNR is then used as the corruption level to generate the noisy version of the utterance. The training process of the DAE model is the same as the one used to train the DNN acoustic model, except that the objective function is replaced as reconstruction error.

In practice, DAE is used as a particular component of the front-end pipeline. The input involves a window of 11 frames Fbank features (after mean normalization), and the output is the noise-removed feature corresponding to the central frame. The outputs are then treated as regular Fbank features that are spliced, LDA transformed and globally normalized before fed into the DNN acoustic model (trained with clean speech).

The CER/PER results of the DNN baseline system on noisy data with the DAE-based noise cancellation are reported in Table 4. Compared to the numbers in Table 3, it can be observed that the performance is significantly improved after noise removed by DAE. This approach is flexible to address any noise by learning a noise-specific DAE, only with a small piece of noise example.

Table 4: Performance of THCHS-30 baseline system with DAE-based noise cancelation.

Condition	Word Task (CER%)	Phone Task (PER%)
0db–White	79.89	51.93
0db–Car	44.91	26.32
0db–Cafeteria	64.75	40.34

4 Conclusions

The aim of the paper is to publish a free 30-hour Chinese speech database, THCHS-30. Additional resources such as lexica, LMs and training recipes are also published to assistant new researchers building their first ASR systems. We demonstrated the building process and presented the baseline results on both clean and noisy data.

As far as we know, this is the first free Chinese speech database that can be used to build a practical Chinese ASR system. We hope this release will attract more potential researchers to this promising field, particularly young researchers whose initial interest would otherwise killed by the high licence fee of commercial databases. We call for challenges based on this database, and hope that invokes deep innovation and more collaboration.

References

1. Li Deng, Jinyu Li, Jui-Ting Huang, Kaisheng Yao, Dong Yu, Frank Seide, Michael L. Seltzer, Geoff Zweig, Xiaodong He, Jason Williams, Yifan Gong, and Alex Acero, "Recent advances in deep learning for speech research at microsoft," *Acoustics, Speech and Signal Processing (ICASSP), 2013 IEEE International Conference on*, 2013.

2. Xuedong Huang, James Baker, and Raj Reddy, "A historical perspective of speech recognition," *Communications of the ACM*, vol. 57, no. 1, pp. 94–103, 2014.

3. Dong Yu and Li Deng, *Automatic Speech Recognition A Deep Learning Approach*, Springer, 2015.

4. Dong Yu and Li Deng, *Deep Learning Methods and Applications*, Foundations and Trends in Signal Processing, 2014.

5. Chenghao Cai, Yanyan Xu, Dengfeng Ke, and Kaile Su, "A fast learning method for multilayer perceptrons in automatic speech recognition systems," *Journal of Robotics*, vol. 2015, no. 2015, pp. 1–7, 2015.

6. Xiaodong He, Li Deng, and Wu Chou, "Discriminative learning in sequential pattern recognition," *IEEE SIGNAL PROCESSING MAGAZINE*, pp. 14–36, 2008.

7. Carla Lopes and Fernando Perdigo, *Phoneme Recognition on the TIMIT Database*, Speech Technologies, 2011.

8. D. Paul and J. Baker, "The design of wall street journal-based csr corpus," *Proceedings of the International Conference on Spoken Language Systems (ICSLP)*, pp. 899–902, 1992.

9. J.J. Godfrey, E.C. Holliman, and J. McDaniel, "Switchboard: telephone speech corpus for research and development," in *Acoustics, Speech, and Signal Processing, 1992. ICASSP-92., 1992 IEEE International Conference on*, Mar 1992, vol. 1, pp. 517–520 vol.1.

10. Thomas Hain, Vincent Wan, Lukas Burget, Martin Karafiat, John Dines, Jithendra Vepa, Giulia Garau, and Mike Lincoln, "The ami system for the transcription of speech in meetings," in *Acoustics, Speech and Signal Processing, 2007. ICASSP 2007. IEEE International Conference on*. IEEE, 2007, vol. 4, pp. IV–357.

11. Aijun Li, Zhigang Yin, Tianqing Wang, Qiang Fang, and Fang Hu, "Rasc863-a chinese speech corpus with four regional accents," *ICSLT-o-COCOSDA, New Delhi, India*, 2004.

12. Askar Rouze, Shi Yin, Zhiyong Zhang, Dong Wang, Askar Humdulla, and Fang Zheng, "Thugy20: A free uyghur speech database," in *NCMMSC 2015*, 2015.

13. Dong Wang, Dalei Wu, and Xiaoyan Zhu, "TCMSD: a new chinese continuous speech database," in *International Conference on Chinese Computing (ICCC'01), 2001,*, 2001.

14. Andreas Stolcke, "SRILM-an extensible language modeling toolkit," in *ICSLP2002*, 2002.

15. Daniel Povey, Arnab Ghoshal, Gilles Boulianne, Lukas Burget, Ondrej Glembek, Nagendra Goel, Mirko Hannemann, Petr Motlicek, Yanmin Qian, Petr Schwarz, Jan Silovsky, Georg Stemmer, and Karel Vesely, "The kaldi speech recognition toolkit," in *IEEE 2011 Workshop on Automatic Speech Recognition and Understanding*. Dec. 2011, IEEE Signal Processing Society, IEEE Catalog No.: CFP11SRW-USB.

16. K. Vesely, A. Ghosal, L. Burget, and D. Povey, "Sequence-discriminative training of deep neural networks,"

Proceedings of the 14th Annual Conference of the International Speech Communication Association, 2013.

17. Dong Yu, Michael L Seltzer, Jinyu Li, Jui-Ting Huang, and Frank Seide, "Feature learning in deep neural networks - a study on speech recognition tasks," in *Proc. of International Conference on Learning Representations*, 2013.

18. Xiangtao Meng, Chao Liu, Zhiyong Zhang, and Dong Wang, "Noisy training for deep neural networks," in *Proc. of ChinaSIP 2014*, 2014, pp. 16–20.

19. Shi Yin, Chao Liu, Zhiyong Zhang, Yiye Lin, Dong Wang, Javier Tejedor, Thomas Fang Zheng, and Yinguo Li, "Noisy training for deep neural networks in speech recognition," *EURASIP Journal on Audio, Speech, and Music Processing*, vol. 2015, no. 1, pp. 1–14, 2015.

20. Pascal Vincent, Hugo Larochelle, Yoshua Bengio, and Pierre-Antoine Manzagol, "Extracting and composing robust features with denoising autoencoders," in *Proceedings of the 25th international conference on Machine learning*. ACM, 2008, pp. 1096–1103.

21. Andrew L Maas, Quoc V Le, Tyler M O'Neil, Oriol Vinyals, Patrick Nguyen, and Andrew Y Ng, "Recurrent neural networks for noise reduction in robust asr.," in *INTERSPEECH*. Citeseer, 2012.

22. Takaaki Ishii, Hiroki Komiyama, Takahiro Shinozaki, Yasuo Horiuchi, and Shingo Kuroiwa, "Reverberant speech recognition based on denoising autoencoder.," in *INTERSPEECH*, 2013, pp. 3512–3516.

23. Mengyuan Zhao, Dong Wang, Zhiyong Zhang, and Xuewei Zhang, "Music removal by denoising autoencoder in speech recognition," in *APSIPA2015*, 2015.